

# **Mastering Artificial Intelligence & Machine Learning**

- From Foundations to ML, DL, NLP, GenAI, Agentic AI & MLOps.

**Author  
Srinivas Dande**





## Course Overview:

- ✓ This course provides a comprehensive, industry-aligned learning path for mastering Artificial Intelligence and Machine Learning. Designed with real-world applications and modern best practices, the program equips learners with the skills required to build, deploy, and optimize AI/ML models.
- ✓ Whether you are aiming for a career in AI, Data Science, or Machine Learning Engineering, this course offers a structured, hands-on approach to mastering AI concepts — from core fundamentals to advanced techniques.
- ✓ **This course is taught step-by-step. No prior AI/ML experience is required.**

## Who Should Attend:

- ✓ **Software Engineers & Developers**
  - Professionals looking to transition into AI/ML, Data Science, or AI Engineering roles.
- ✓ **Data Scientists & ML Engineers**
  - Practitioners who want to upskill with advanced AI/ML, Generative AI, LLMs, Agentic AI, and MLOps concepts.
- ✓ **Tech Professionals & AI Enthusiasts**
  - Individuals aiming to break into the AI/ML industry with a structured, hands-on learning approach.
- ✓ **Professionals Preparing for AI/ML Interviews**
  - Candidates preparing for AI/ML roles at top product-based companies through strong fundamentals, projects, and interview-focused preparation.
- ✓ **Job Seekers Targeting High-Growth AI/ML Roles**
  - Learners aiming for high-impact and high-growth AI/ML roles in leading technology-driven organizations.

## Why Dandes Academy?

- ✓ Choosing the right AI/ML program is not just about learning algorithms — it's about building real, industry-ready skills that translate into long-term career growth.
- ✓ At Dandes Academy, we focus on depth, structure, and real-world application, not shortcuts.

### 1) Industry-Aligned, Future-Ready Curriculum

- Our 16-module AI/ML program is carefully designed to take you from foundations to advanced, production-level AI systems.
- You don't just learn models — you learn how to:
  - ✓ Build end-to-end AI/ML solutions
  - ✓ Work with Generative AI, LLMs, Agentic AI
  - ✓ Deploy and manage models using MLOps best practices
- This ensures you're ready not just for today's roles, but for future AI careers.

### 2) Step-by-Step Learning (Beginner Friendly)

- This course is taught step-by-step.
  - ✓ No prior AI/ML experience is required.
- **We start with:**
  - ✓ Python, SQL, Data Modeling
  - ✓ Mathematics, Statistics, EDA
  - ✓ And gradually move into:
  - ✓ Machine Learning → Deep Learning
  - ✓ NLP → Generative AI → LLMs
  - ✓ Agentic AI → MLOps → Capstone Project
- This structured approach ensures clarity, confidence, and strong fundamentals.



## 3) Founder-Led Teaching with 20+ Years of Experience

- This program is taught by **Srinivas Dande, Founder & Instructor at Dandes Academy**, with **20+ years of industry and teaching experience**.
- **What this means for you:**
  - ✓ Clear explanations from basics to advanced
  - ✓ Real-world insights, not just theory
  - ✓ Guidance aligned with actual industry expectations

## 4) Hands-On Learning with Real Projects

- Learning happens by doing, not just watching.
- **You will work on:**
  - ✓ Multiple mini-projects
  - ✓ Real-time industry-oriented projects
  - ✓ A final end-to-end AI/ML capstone project
- This builds a strong project portfolio that recruiters value.

## 5) Advanced Skills That Most Courses Don't Teach

- **Unlike many programs, we go deep into:**
  - ✓ Generative AI & LLMs
  - ✓ Fine-tuning & RAG pipelines
  - ✓ Agentic AI (Autonomous & Multi-Agent Systems)
  - ✓ Production MLOps & deployment
- These are the skills companies are actively looking for right now.

## 6) Strong Career & Placement Focus

- We don't just teach — we prepare you for the job market.
- **You get:**
  - ✓ Placement assistance
  - ✓ Resume & interview guidance
  - ✓ Mock interviews
  - ✓ Internship opportunities (based on performance)
- This helps you confidently target roles such as:
  - ✓ ML Engineer,
  - ✓ AI Engineer,
  - ✓ Gen AI Engineer,
  - ✓ Data Scientist,
  - ✓ Lead ML Engineer,
  - ✓ ML Architect.

## 7) Long-Term Learning & Flexibility

- 12 months structured program
- 5 years LMS access for recordings & materials
- Batch switching option (subject to availability)
- 100% refund within 30 days if you're not satisfied — no questions asked



## What You Will Learn:

- ✓ **Strong AI & ML Foundations**
  - Master core concepts in Supervised and Unsupervised Learning, along with Statistics, Mathematics for ML, and Exploratory Data Analysis (EDA) to build a solid foundation.
- ✓ **End-to-End Machine Learning**
  - Design, train, evaluate, and optimize Machine Learning models using industry-standard techniques, including feature engineering, model selection, and hyperparameter tuning.
- ✓ **Deep Learning & Neural Networks**
  - Gain hands-on experience with Deep Learning architectures, including ANNs, CNNs, and RNNs, and apply them to real-world problems.
- ✓ **Natural Language Processing (NLP)**
  - Work with text data, covering text preprocessing, embeddings, transformers, and NLP use cases.
- ✓ **Generative AI & Large Language Models (LLMs)**
  - Explore Generative AI, including Transformers, Diffusion Models, LLMs, and fine-tuning techniques, and build practical GenAI applications.
- ✓ **Agentic AI Systems**
  - Learn to design and build Agentic AI workflows, including multi-step reasoning, tool usage, and real-world agent-based applications.
- ✓ **MLOps & Model Deployment**
  - Build end-to-end AI pipelines using MLOps best practices, including model deployment, versioning, monitoring, and automation.
- ✓ **Hands-On Projects & Real-Time Case Studies**
  - Work on multiple mini-projects and real-time industry projects to gain practical exposure and build a strong project portfolio.
- ✓ **Industry-Ready AI Skills**
  - Gain practical skills to design, deploy, and scale AI/ML solutions, preparing you for real-world roles such as ML Engineer, AI Engineer, and ML Architect.



## List of Modules

Course 1: Python Fundamentals

Course 2: Python for ML & DS

Course 3: SQL, Advanced SQL and Data Modeling

Course 4: Mathematics for ML & DS

Course 5: Exploratory Data Analysis (EDA)

Course 6: Machine Learning - Foundation

Course 7: Advanced Machine Learning

Course 8: Deep Learning - Foundation

Course 9: Advanced Deep Learning

Course 10: Image Processing & Computer Vision

Course 11: Natural Language Processing (NLP)

Course 12: Generative AI

Course 13: Large Language Models (LLMs) & Fine-Tuning

Course 14: Agentic AI (NEW)

Course 15: MLOps

Course 16: End-to-End AI/ML Project

## Course 1: Python Fundamentals

- **Topic Coverage**

- ✓ Data Types & Variables
- ✓ Conditional Statements (if-else)
- ✓ Loops (for, while)
- ✓ Functions & Lambda Functions
- ✓ Strings, Lists, Tuples, Sets, Dictionaries
- ✓ Comprehensions (List, Dict, Set)
- ✓ Functional Programming
- ✓ Object-Oriented Programming
- ✓ Decorators
- ✓ Generators & Iterators
- ✓ Exception Handling
- ✓ File Handling

- **Assessment**

- ✓ Practice Questions
- ✓ Assignment (Python Problem Solving)
- ✓ Test 1 - Python Fundamentals

## Course 2: Python for ML & DS

- **Topic Coverage**

- ✓ NumPy
- ✓ Pandas
- ✓ Matplotlib
- ✓ Seaborn

- **Assessment**

- ✓ Practice Questions
- ✓ Assignment #1 - Pandas
- ✓ Assignment #2 - Plotting
- ✓ Test 2 - Python for ML & DS

## Course 3: SQL, Advanced SQL, and Data Modeling

- **Topic Coverage**

- ✓ SQL Query Fundamentals
- ✓ Set Operators
- ✓ Group By & Aggregations
- ✓ Joins
- ✓ Common Table Expressions (CTEs)
- ✓ Window Functions
- ✓ RANK, ROW\_NUMBER, LEAD, LAG, etc
- ✓ Nested Queries
- ✓ Indexes
- ✓ Query Optimization & Best Practices
- ✓ Data Modeling (OLTP vs OLAP)
- ✓ Entity Relationships
- ✓ Normalization
- ✓ Star Schema & Snowflake Schema

- **Assessment**

- ✓ Practice Questions
- ✓ Assignment #1 - Advanced SQL
- ✓ Assignment #2 - Data Modeling
- ✓ Test 3 - SQL, Advanced SQL, and Data Modeling

## Course 4: Mathematics for ML & DS

- **Topic Coverage**

- ✓ Probability Basics & Rules
- ✓ Probability Distributions
- ✓ Important Theorems
- ✓ Descriptive & Inferential Statistics
- ✓ Hypothesis Testing
- ✓ Correlation & Covariance
- ✓ Linear Algebra
- ✓ Vectors, Matrices
- ✓ Eigenvalues & Eigenvectors
- ✓ Matrix Decomposition
- ✓ Calculus for ML
- ✓ Derivatives & Gradients
- ✓ Optimization Intuition
- ✓ Regression Mathematics (Linear & Polynomial)

- **Assessment**

- ✓ Practice Questions
- ✓ Assignment (Math → ML Mapping)
- ✓ Test 4 - Mathematics for ML & DS

## Course 5: Exploratory Data Analysis (EDA)

- **Topic Coverage**

- ✓ Understanding the Data
- ✓ Data Cleaning & Validation
- ✓ Handling Missing Data & Outliers
- ✓ Feature Engineering
- ✓ \* Encoding Techniques
- ✓ \* Scaling & Normalization
- ✓ \* Feature Transformation
- ✓ \* Time-based Features
- ✓ Dimensionality Reduction (PCA, VIF)
- ✓ Data Leakage Detection
- ✓ Model Evaluation Metrics
- ✓ Train / Validation / Test Split
- ✓ Stratified Sampling
- ✓ Data-Centric AI Principles

- **Assessment**

- ✓ Practice Questions
- ✓ Assignment (EDA + Feature Engineering Project)
- ✓ Test 5 - Exploratory Data Analysis (EDA)

## Course 6: Machine Learning - Foundation

- **Topic Coverage**

- ✓ Introduction to Machine Learning
- ✓ Supervised vs Unsupervised Learning
- ✓ Supervised Learning
  - ✓ Regression Algorithms
    - Linear Regression
    - Regularization (L1, L2)
- ✓ Classification Algorithms
  - Logistic Regression
  - Naïve Bayes
  - K-Nearest Neighbors (KNN)
  - Support Vector Machines (SVM)
  - Handling Imbalanced Data
- ✓ Unsupervised Learning
- ✓ Clustering Techniques
  - K-Means
  - Hierarchical Clustering
  - Elbow Method
  - Clustering Metrics & Evaluation
- ✓ Evaluation Metrics
  - Regression Metrics
  - Classification Metrics

- **Assessment**

- ✓ Practice Questions
- ✓ Assignment - Supervised Learning
- ✓ Assignment - Unsupervised Learning
- ✓ Test 6 - Machine Learning - Foundation



## Course 7: Advanced Machine Learning

- **Topic Coverage**

- ✓ Advanced Supervised Learning
- ✓ Decision Trees
- ✓ Bagging Techniques
- ✓ Random Forest
- ✓ Boosting Techniques
  - AdaBoost
  - Gradient Boosting
  - XGBoost
- ✓ Advanced Unsupervised Learning
- ✓ Advanced Clustering
  - DBSCAN
  - Birch
  - GMM
- ✓ Expectation Maximization
- ✓ Anomaly Detection
- ✓ Dimensionality Reduction
- ✓ Model Selection & Optimization
- ✓ Time Series Forecasting

- **Assessment**

- ✓ Practice Questions
- ✓ Mini Project #1 - Supervised Learning
- ✓ Mini Project #2 - Unsupervised Learning
- ✓ Test 7 - Advanced Machine Learning

## Course 8: Deep Learning - Foundation

- **Topic Coverage**

- ✓ Deep Learning Fundamentals
- ✓ Basics of Neural Networks
- ✓ Forward & Backpropagation
- ✓ Gradient Descent & Optimization
- ✓ Activation Functions
  - Sigmoid, tanH, ReLU
  - Leaky ReLU, Softmax
- ✓ Training Deep Neural Networks
- ✓ Weight Initialization Strategies
- ✓ Regularization Techniques
  - L1, L2, Dropout
  - Batch Normalization
- ✓ Learning Rate Schedulers
- ✓ Optimizers - SGD, Adam, RMSprop-
- ✓ Tensors & Computation
- ✓ Deep Learning Frameworks
  - PyTorch and TensorFlow
- ✓ Core Deep Learning Architectures
  - Convolutional Neural Networks (CNNs)
  - Recurrent Neural Networks (RNNs)
  - Long Short-Term Memory (LSTMs)
- ✓ Hyperparameter Tuning & Model Optimization

- **Assessment**

- ✓ Practice Questions
- ✓ Assignment
- ✓ Test 8 - Deep Learning - Foundation

## Course 9: Advanced Deep Learning

- **Topic Coverage**

- ✓ Advanced Deep Learning Architectures
  - Advanced CNN Architectures
  - RNN, LSTM, GRU
  - Transfer Learning
- ✓ Advanced Training & Optimization
  - Deep Learning Model Optimization
  - Training Stability Issues
  - Cross-Validation in Deep Learning
  - Hyperparameter Search
- ✓ Model Explainability & Responsible AI
  - Model Interpretability - SHAP and LIME
  - AI Fairness & Bias Reduction
- ✓ Transformers
  - Introduction to Transformers
  - High-Level Overview of Transformer Architectures

- **Assessment**

- ✓ Practice Questions
- ✓ Mini Project #3 : Deep Learning Applications
- ✓ Test 9: Advanced Deep Learning

## Course 10: Computer Vision

- **Topic Coverage**

- ✓ Image Processing Fundamentals
- ✓ Convolutional Neural Networks for Vision
- ✓ Fundamentals of CNNs
- ✓ Modern CNN Architectures
  - SqueezeNet
  - ResNet
  - EfficientNet
- ✓ Image Classification
  - Image Classification Concepts
  - Pre-trained CNN Models
  - Custom CNN Models
- ✓ Object Detection
  - Object Detection Fundamentals
  - YOLO, SSD, Faster R-CNN
- ✓ Image Augmentation & Preprocessing
  - Image Augmentation Techniques
  - Data Preprocessing for Vision

- **Assessment**

- ✓ Practice Questions
- ✓ Mini Project #4 : Image Classification
- ✓ Mini Project #5: Real-Time Object Detection
- ✓ Test 10: Computer Vision

## Course 11: Natural Language Processing (NLP)

- **Topic Coverage**

- ✓ NLP Fundamentals
  - Introduction to NLP
  - Text Preprocessing & Tokenization
- ✓ Text Representation
  - Word Embeddings
  - Word2Vec, GloVe, FastText
- ✓ Deep Learning for NLP
  - \* Neural Networks for NLP
  - \* Seq2Seq Models
  - \* Transformers for NLP
- ✓ Pretrained Language Models
  - BERT
  - T5
  - Other Pretrained Models
- ✓ NLP Applications
  - Named Entity Recognition (NER)
  - Sentiment Analysis
  - Sentiment Analysis on Product Reviews

- **Assessment**

- ✓ Practice Questions
- ✓ Mini Project #6 : AI Chatbot for Customer Support
- ✓ Test 11: Natural Language Processing (NLP)

## Course 12: Generative AI

- **Topic Coverage**

- ✓ Introduction to Generative AI
- ✓ Applications in Text, Image, and Video Generation
- ✓ Autoencoders
  - Introduction to Autoencoders
  - Variational Autoencoders (VAEs)
  - Denoising Autoencoders
  - Sparse Autoencoders
- ✓ Generative Adversarial Networks (GANs)
  - GAN Fundamentals
  - GAN Variants
- ✓ Diffusion Models
  - Diffusion Models Overview
  - Stable Diffusion
  - DALL·E
  - MidJourney
- ✓ Generative AI Applications
  - AI-Generated Art using GANs
  - Image-to-Image Translation with CycleGAN
  - Text-to-Image Generation with Stable Diffusion

- **Assessment**

- ✓ Practice Questions
- ✓ Mini Project #7 : AI-Generated Art & Image-to-Image Translation
- ✓ Test 12: Generative AI



## Course 13: LLMs & Fine-Tuning

- **Topic Coverage**

- ✓ Introduction to Large Language Models
- ✓ Evolution of LLMs- GPT, BERT, T5, LLaMA, Claude, Gemini
- ✓ LLM Foundations
  - Tokenization Strategies
  - WordPiece, SentencePiece, Byte Pair Encoding (BPE)
  - Pretraining vs Fine-Tuning vs Prompt Engineering
- ✓ LLM Frameworks & Tooling
  - LangChain and LlamaIndex
- ✓ Fine-Tuning Large Language Models
  - Fine-Tuning Open-Source LLMs - LLaMA, Falcon
  - Parameter-Efficient Fine-Tuning (PEFT) and LoRA
- ✓ Alignment & Optimization
  - Reinforcement Learning with Human Feedback (RLHF)
  - Reward Models and Multi-Armed Bandits
  - Proximal Policy Optimization (PPO) for LLMs
- ✓ Retrieval-Augmented Generation (RAG)
  - RAG Architecture & Workflow
  - Building RAG Pipelines for LLMs
- ✓ LLM Applications
  - Building AI Chatbots & Assistants
  - Code Generation using LLMs- CodeLlama, StarCoder
  - Text Summarization and AI-Powered Content Creation

- **Assessment**

- ✓ Practice Questions
- ✓ Mini Project #8 : Fine-Tuning an LLM for a Domain-Specific Chatbot
- ✓ Test 13: LLMs & Fine-Tuning

## Course 14: Agentic AI (Autonomous & Multi-Agent Systems)

- **Topic Coverage**

- ✓ Introduction to Agentic AI
  - What is Agentic AI
  - LLMs as Reasoning Agents
  - Agentic AI vs Traditional AI Pipelines
  - Single-Agent vs Multi-Agent Systems
- ✓ Core Agent Concepts
  - Agent Architecture
  - Perception, Planning, Reasoning, Action Loop
  - Tool Calling & Function Calling
  - Prompting Strategies for Agents
- ✓ Memory in Agents
  - Short-Term Memory
  - Long-Term Memory
  - Conversation Memory
  - Vector Memory using Embeddings
- ✓ Planning & Reasoning
  - Task Decomposition
  - Planning Strategies (ReAct, Plan-and-Execute)
  - Chain-of-Thought & Reasoning Patterns
  - Handling Failures & Retries
- ✓ Agent Frameworks & Tooling
  - LangChain Agents
  - LlamaIndex Agents
  - Tool Integration (APIs, Databases, Web Tools)
- ✓ Multi-Agent Systems
  - Multi-Agent Architectures
  - Agent Collaboration & Coordination
  - Role-Based Agents
  - Task Delegation among Agents



- ✓ Autonomous AI Agents
  - Autonomous Decision Making
  - Long-Running Agents
  - Event-Driven Agents
  - Human-in-the-Loop Systems
- ✓ Agent Evaluation, Safety & Control
  - Agent Evaluation Metrics
  - Guardrails & Safety Mechanisms
  - Preventing Hallucinations
  - Cost & Latency Optimization
  - Ethical Considerations in Agentic AI
- ✓ Agentic AI Applications
  - AI Personal Assistants
  - Autonomous Research Agents
  - Customer Support AI Agents
  - Data Analysis & Reporting Agents

- **Assessment**

- ✓ Practice Questions
- ✓ Mini Project #9: Build an Autonomous AI Agent
- ✓ Test 14: Agentic AI

## Course 15: MLOps

- **Topic Coverage**
  - ✓ ML Lifecycle & Foundations
    - End-to-End ML Model Lifecycle
    - MLOps Best Practices
    - Experiment Tracking & Reproducibility
  - ✓ Model & Feature Management
    - Model Versioning
    - Feature Store Management
    - Data Versioning
  - ✓ MLOps Tooling & Platforms
    - MLflow, Kubeflow
    - TensorFlow Extended (TFX)
    - Kubernetes for Machine Learning
  - ✓ CI/CD for Machine Learning
    - CI/CD Pipelines for ML
    - GitOps for ML
    - Kubeflow Pipelines
    - Automated Training & Deployment
  - ✓ Model Deployment
    - Model Serving & Inference
    - Online vs Batch Inference
  - ✓ Monitoring, Evaluation & Optimization
    - Model Monitoring
    - Data Drift vs Concept Drift
    - Model Performance Tracking
- **Assessment**
  - ✓ Practice Questions
  - ✓ Assignment
  - ✓ Test 15: MLOps

## Course 16: Final End-to-End AI/ML Project

- **Topic Coverage**

- ✓ **Capstone Project**

- AI-Powered Personal Assistant (Multi-Modal AI Agent)

- ✓ **Project Scoping & Design**

- Problem Definition
    - Requirements Gathering
    - System Architecture Design

- ✓ **Data Collection & Preprocessing**

- Data Collection, Cleaning & Transformation
    - Dataset Versioning

- ✓ **Feature Engineering & Selection**

- Feature Engineering
    - Feature Selection Techniques

- ✓ **Model Training & Optimization**

- Model Selection
    - Training Pipelines
    - Hyperparameter Optimization

- ✓ **Model Evaluation & Validation**

- Evaluation Metrics
    - Validation Strategies and Error Analysis

- ✓ **Model Deployment**

- API Deployment
    - Web Application Deployment

- ✓ **Model Monitoring & Maintenance**

- Model Monitoring
    - Performance Tracking and Drift Detection

- ✓ **Scaling & Continuous Improvement**

- Scalability Strategies and Cost Optimization
    - Continuous Model Improvement